

Problem Solutions

Chapter 3

Pierre Paquay

Problem 3.1


First, we generate 2000 examples uniformly for the two semi-circles, this means that we will have approximately 1000 examples for each class.

```
set.seed(101)

init_data <- function(N, rad, thk, sep) {
  D <- data.frame(x = numeric(), y = numeric())
  y <- numeric()
  repeat {
    x1 <- runif(1, min = -25, max = 40)
    x2 <- runif(1, min = -30, max = 20)
    if ((x2 >= 0) && (rad^2 <= x1^2 + x2^2) && (x1^2 + x2^2 <= (rad + thk)^2)) {
      D <- rbind(D, c(x1, x2))
      y <- c(y, -1)
    }
    else if ((x2 < -sep) && (rad^2 <= (x1 - rad - thk / 2)^2 + (x2 + sep)^2) &&
      ((x1 - rad - thk / 2)^2 + (x2 + sep)^2 <= (rad + thk)^2)) {
      D <- rbind(D, c(x1, x2))
      y <- c(y, +1)
    }
    if (nrow(D) >= N)
      break
  }
  colnames(D) <- c("x1", "x2")

  return(cbind(D, y))
}

rad <- 10
thk <- 8
sep <- 5
D <- init_data(2000, rad, thk, sep)
p <- ggplot(D, aes(x = x1, y = x2, col = as.factor(y + 3))) + geom_point() +
  theme(legend.position = "none") +
  coord_fixed()
p
```




(a) Then, we run the PLA starting from $w = 0$ until it converges and we plot the data with the final hypothesis.

```
h <- function(D, w) {
  scalar_prod <- cbind(1, D$x1, D$x2) %*% w


  return(as.vector(sign(scalar_prod)))
}

iter <- 0
w_PLA <- c(0, 0, 0)
repeat {
  y_pred <- h(D, w_PLA)
  D_mis <- subset(D, y != y_pred)
  if (nrow(D_mis) == 0)
    break
  xt <- D_mis[1, ]
  w_PLA <- w_PLA + c(1, xt$x1, xt$x2) * xt$y
  iter <- iter + 1
}
p + geom_abline(slope = -w_PLA[2] / w_PLA[3], intercept = -w_PLA[1] / w_PLA[3])
```



(b) Now, we use linear regression for classification to obtain w_{lin} .

```
X <- as.matrix(cbind(1, D[, c("x1", "x2")]))
y <- D$y
X_cross <- solve(t(X) %*% X) %*% t(X)
w_lin <- as.vector(X_cross %*% y)
p + geom_abline(slope = -w_lin[2] / w_lin[3], intercept = -w_lin[1] / w_lin[3])
```




As we may see, linear regression can also be used for classification (the values $\text{sign}(w_{lin}^T x)$ will likely make good classification predictions). The linear regression weights w_{lin} are also an approximate solution for the perceptron model.

Problem 3.2

In this problem, we consider again the double-semi-circle of Problem 3.1 and we vary sep in the range $\{0.2, 0.4, \dots, 5\}$, with these values we generate 2000 examples and we run PLA starting with $w = 0$. Below, we plot sep versus the number of iterations PLA takes to converge.

```
set.seed(10)


sep_seq <- seq(0.2, 5, 0.2)
iterations <- numeric()
iterations_max <- numeric()
for (sep in sep_seq) {
  D <- init_data(2000, rad, thk, sep)
  iter <- 0
  w_PLA <- c(0, 0, 0)
  repeat {
    y_pred <- h(D, w_PLA)
    D_mis <- subset(D, y != y_pred)
    if (nrow(D_mis) == 0)
      break
    xt <- D_mis[1, ]
    w_PLA <- w_PLA + c(1, xt$x1, xt$x2) * xt$y
    iter <- iter + 1
  }
  iterations <- c(iterations, iter)
  R <- max(apply(cbind(1, D[, 1:2]), 1, FUN = function(x) sqrt(1 + x[2]^2 + x[3]^2)))
  Rho <- min(D$y * apply(cbind(1, D[, 1:2]), 1, FUN = function(x) sum(w_PLA * x)))
  iterations_max <- c(iterations_max, R^2 * sum(w_PLA^2) / Rho^2)
}
ggplot(data.frame(sep = sep_seq, iterations = iterations), aes(x = sep, y = iterations)) +
  geom_line()
```



We may see that the number of iterations tends to decrease when sep increases. This trend is confirmed

by the theoretical results (see Problem 1.3), to see this we plot below *sep* versus the maximum number of iterations.


```
ggplot(data.frame(sep = sep_seq, iterations = iterations_max), aes(x = sep, y = iterations)) +  
  geom_line(col = "red") +  
  ylab("Maximum number of iterations (theoretical)")
```



Problem 3.3

Here again, we consider the double-semi-circle of Problem 3.1 and we set $sep = -5$ (which makes the data non linearly separable) and we generate 2000 examples.

```
set.seed(101)  
  
sep <- -5  
D <- init_data(2000, rad, thk, sep)  
p <- ggplot(D, aes(x = x1, y = x2, col = as.factor(y + 3))) + geom_point() +  
  theme(legend.position = "none") +  
  coord_fixed()  
p
```




(a) If we run PLA on these examples, it will never stop updating.

(b) Now, we run the pocket algorithm for 100000 iterations and we plot E_{in} versus the iteration number t for $t = 1, \dots, 100$.

```
set.seed(101)


start_time_pocket <- Sys.time()
E_in <- numeric()
E_in_pocket <- numeric()
w <- c(0, 0, 0)
w_pocket <- w
E_in <- c(E_in, mean(D$y != h(D, w)))
E_in_pocket <- E_in
for (iter in 1:100000) {
  D_mis <- subset(D, y != h(D, w))
  if (nrow(D_mis) == 0)
    break
  xt <- D_mis[sample(nrow(D_mis), 1), ]
  w <- w + c(1, xt$x1, xt$x2) * xt$y
  E_in <- c(E_in, mean(D$y != h(D, w)))
  if (E_in[length(E_in)] < E_in_pocket[length(E_in_pocket)]) {
    w_pocket <- w
  }
  E_in_pocket <- c(E_in_pocket, mean(D$y != h(D, w_pocket)))
}
end_time_pocket <- Sys.time()
ggplot(data.frame(t = 1:100000, E_in = E_in_pocket[-1]), aes(x = t, y = E_in)) +
  geom_line(col = "red") +
  coord_cartesian(xlim = c(1, 100))
```



We clearly see that the E_{in} is monotonously decreasing (as opposed to what would happen if we used the PLA).

(c) Below, we plot the data and the final hypothesis obtained in (b).

```
p + geom_abline(slope = -w_pocket[2] / w_pocket[3], intercept = -w_pocket[1] / w_pocket[3])
```




(d) Here, we use the linear regression algorithm to obtain w and we compare this result with the pocket algorithm in terms of computation time and quality of the solution.

```
start_time_lin <- Sys.time()
X <- as.matrix(cbind(1, D[, c("x1", "x2")]))
y <- D$y
X_cross <- solve(t(X) %*% X) %*% t(X)
```

```

w_lin <- as.vector(X_cross %*% y)
E_in_lin <- mean(D$y != h(D, w_lin))
end_time_lin <- Sys.time()
p + geom_abline(slope = -w_lin[2] / w_lin[3], intercept = -w_lin[1] / w_lin[3])

```



For the pocket algorithm (with 100000 iterations), we have a computation time of 1.7378043, and for the linear regression algorithm, we have a computation time of 0.0046861. The linear regression algorithm is clearly better than the pocket algorithm in terms of computation time. When we take into account the quality of the solution, the pocket algorithm has a (final) E_{in} of 0.086, and the linear regression algorithm has a E_{in} of 0.0995. So, regarding the quality of the solution, the pocket algorithm is a little better than the linear regression algorithm.

(e) Here, we repeat the points (b) to (d) with a 3rd order polynomial feature transform

$$\Phi(x) = (1, x_1, x_2, x_1^2, x_1x_2, x_2^2, x_1^3, x_1^2x_2, x_1x_2^2, x_2^3).$$

First, we run the pocket algorithm for 100000 iterations and we plot E_{in} versus the iteration number t .

```

set.seed(11)

D_trans <- data.frame(x1 = D$x1, x2 = D$x2,
  x1_sq = D$x1^2, x1x2 = D$x1 * D$x2, x2_sq = D$x2^2,
  x1_cub = D$x1^3, x1_sqx2 = D$x1^2 * D$x2,
  x1x2_sq = D$x1 * D$x2^2, x2_cub = D$x2^3, y = D$y)

h_trans <- function(D, w) {
  scalar_prod <- as.matrix(cbind(1, D[, 1:9])) %*% w

  return(as.vector(sign(scalar_prod)))
}


```



```

start_time_pocket <- Sys.time()
E_in <- numeric()
E_in_pocket <- numeric()
w <- rep(0, 10)
w_pocket <- w
E_in <- c(E_in, mean(D_trans$y != h_trans(D_trans, w)))
E_in_pocket <- E_in
for (iter in 1:100000) {
  D_mis <- subset(D_trans, y != h_trans(D_trans, w))
  if (nrow(D_mis) == 0)
    break
  xt <- D_mis[sample(nrow(D_mis), 1), ]
  w <- w + c(1, as.numeric(xt[1:9])) * xt$y
  E_in <- c(E_in, mean(D_trans$y != h_trans(D_trans, w)))
  if (E_in[length(E_in)] < E_in_pocket[length(E_in_pocket)]) {
    w_pocket <- w
  }
  E_in_pocket <- c(E_in_pocket, mean(D_trans$y != h_trans(D_trans, w_pocket)))
}
end_time_pocket <- Sys.time()
ggplot(data.frame(t = 1:100000, E_in = E_in_pocket[-1]), aes(x = t, y = E_in)) +
  geom_line(col = "red") +
  coord_cartesian(xlim = c(1, 100))

```




Then, we plot the data and the final hypothesis obtained above.

```


cc <- emdbook::curve3d(1 * w_pocket[1] + x * w_pocket[2] + y * w_pocket[3] + x^2 * w_pocket[4] +
  x * y * w_pocket[5] + y^2 * w_pocket[6] + x^3 * w_pocket[7] +
  x^2 * y * w_pocket[8] + x * y^2 * w_pocket[9] + y^3 * w_pocket[10],
  xlim = c(-20, 35), ylim = c(-15, 20), sys3d = "none")
dimnames(cc$z) <- list(cc$x, cc$y)
mm <- reshape2::melt(cc$z)
p + geom_contour(data = mm, aes(x = Var1, y = Var2, z = value), breaks = 0, colour = "black")

```



Finally, we use the linear regression algorithm to obtain the weights w .

```
start_time_lin <- Sys.time()
X <- as.matrix(cbind(1, D_trans[, 1:9]))
y <- D_trans$y
X_cross <- solve(t(X) %*% X) %*% t(X)
w_lin <- as.vector(X_cross %*% y)
E_in_lin <- mean(D_trans$y != h_trans(D_trans, w_lin))
end_time_lin <- Sys.time()
cc <- emdbook::curve3d(1 * w_lin[1] + x * w_lin[2] + y * w_lin[3] + x^2 * w_lin[4] +
  x * y * w_lin[5] + y^2 * w_lin[6] + x^3 * w_lin[7] +
  x^2 * y * w_lin[8] + x * y^2 * w_lin[9] + y^3 * w_lin[10],
  xlim = c(-20, 35), ylim = c(-15, 20), sys3d = "none")
dimnames(cc$z) <- list(cc$x, cc$y)
mm <- reshape2::melt(cc$z)
p + geom_contour(data = mm, aes(x = Var1, y = Var2, z = value), breaks = 0, colour = "black")
```



For the pocket algorithm (with 100000 iterations), we have a computation time of 3.0310437, and for the linear regression algorithm, we have a computation time of 0.00687. The linear regression algorithm is once again clearly better than the pocket algorithm in terms of computation time. When we take into account the quality of the solution, the pocket algorithm has a (final) E_{in} of 0.0445, and the linear regression algorithm has a E_{in} of 0.021. In this case, regarding the quality of the solution, the linear regression algorithm is also better than the pocket algorithm.